

## Assignment 6.1

### Task 1

Use Sqoop tool to export data present in SQOOPOUT folder made while demo of Import table

### Task2

Use Sqoop tool to export data present in SQOOPOUT folder made while demo of Import table with parameter person\_id =3.

### 6. Expected Output

Solution report with commands, explanation to commands and screenshot for output. Report shall be in PDF format. Submitted in GitHub.

## SOLUTION – TASK1

### #dump the database and table name 'Person'

```
mysql> show databases;
+-----+
| Database |
+-----+
| information_schema |
| metastore |
| mysql |
| oozie |
| performance_schema |
| simplidb |
| sys |
| testdb |
+-----+
8 rows in set (0.00 sec)

mysql> use simplidb;
Database changed
mysql> dump tables;
ERROR 1064 (42000): You have an error in your SQL syntax; check the manual that
corresponds to your MySQL server version for the right syntax to use ne
ar 'dump tables' at line 1
mysql>
mysql>
mysql>
mysql> show tables;
+-----+
| Tables_in_simplidb |
+-----+
| Person |
| employee |
+-----+
2 rows in set (0.00 sec)
```

### #Schema of the table Person

```
mysql> describe Person;
```

Field	Type	Null	Key	Default	Extra
person_id	int(11)	NO	PRI	NULL	
lname	varchar(20)	YES		NULL	
fname	varchar(20)	YES		NULL	
area	varchar(20)	YES		NULL	
city	varchar(20)	YES		NULL	

5 rows in set (0.00 sec)

### #Dump contents of the table Person , there are 12 records

```
mysql>
```

```
mysql>
```

```
mysql> select * from Person;
```

person_id	lname	fname	area	city
1	Rizvi	Syed	Whitefield	Bangalore
2	Singh	Amit	Whitefield	Bangalore
3	Singh	Ravi	Whitefield	Bangalore
4	Misra	Raju	Whitefield	Bangalore
5	hulu	bulu	Whitefield	Bangalore
6	Sinha	bulbul	Whitefield	Bangalore
7	Chue	Wang	Whitefield	Bangalore
8	Mehra	Vinod	Whitefield	Bangalore
9	Bachchan	Amitabh	Whitefield	Bombay
10	Bai	Munni	Whitefield	Lucknow
11	Bhai	Laddoo	Koramangala	Lucknow
12	Hosakote	Sri	Koramangala	Lucknow

12 rows in set (0.00 sec)

### #import this table into HDFS using Sqoop

```
[acadgild.mmisra ~]$
```

```
[acadgild.mmisra ~]$ sqoop import --connect jdbc:mysql://localhost/simplidb --table Person --username root -password Root@123 --target-dir /sqoopoutfull -m 1
```

```
Warning: /home/acadgild/install/sqoop/sqoop-1.4.6.bin__hadoop-2.0.4-alpha/./hcatalog does not exist! HCatalog jobs will fail.
```

```
Please set $HCAT_HOME to the root of your HCatalog installation.
```

```
Warning: /home/acadgild/install/sqoop/sqoop-1.4.6.bin__hadoop-2.0.4-alpha/./accumulo does not exist! Accumulo imports will fail.
```

```
Please set $ACCUMULO_HOME to the root of your Accumulo installation.
```

```
18/07/07 12:05:23 INFO sqoop.Sqoop: Running Sqoop version: 1.4.6
```

```
18/07/07 12:05:23 WARN tool.BaseSqoopTool: Setting your password on the command-line is insecure. Consider using -P instead.
```

```
18/07/07 12:05:23 INFO manager.MySQLManager: Preparing to use a MySQL streaming resultset.
```

```
18/07/07 12:05:23 INFO tool.CodeGenTool: Beginning code generation
```

```
Sat Jul 07 12:05:23 IST 2018 WARN: Establishing SSL connection without server's identity verification is not recommended. According to MySQL 5.5.45+, 5.6.26+ and 5.7.6+ requirements SSL connection must be established by default if explicit option isn't set. For compliance with existing applications not using SSL the verifyServerCertificate property is set to 'false'. You need either to explicitly
```

disable SSL by setting useSSL=false, or set useSSL=true and provide truststore for server certificate verification.

18/07/07 12:05:23 INFO manager.SqlManager: Executing SQL statement: SELECT t.\* FROM `Person` AS t LIMIT 1

18/07/07 12:05:24 INFO manager.SqlManager: Executing SQL statement: SELECT t.\* FROM `Person` AS t LIMIT 1

18/07/07 12:05:24 INFO orm.CompilationManager: HADOOP\_MAPRED\_HOME is /home/acadgild/install/hadoop/hadoop-2.6.5

Note: /tmp/sqoop-acadgild/compile/d7876513ad447baf2894d361b74dbc59/Person.java uses or overrides a deprecated API.

Note: Recompile with -Xlint:deprecation for details.

18/07/07 12:05:25 INFO orm.CompilationManager: Writing jar file: /tmp/sqoop-acadgild/compile/d7876513ad447baf2894d361b74dbc59/Person.jar

18/07/07 12:05:25 WARN manager.MySQLManager: It looks like you are importing from mysql.

18/07/07 12:05:25 WARN manager.MySQLManager: This transfer can be faster! Use the --direct

18/07/07 12:05:25 WARN manager.MySQLManager: option to exercise a MySQL-specific fast path.

18/07/07 12:05:25 INFO manager.MySQLManager: Setting zero DATETIME behavior to convertToNull (mysql)

18/07/07 12:05:25 INFO mapreduce.ImportJobBase: Beginning import of Person

SLF4J: Class path contains multiple SLF4J bindings.

SLF4J: Found binding in [jar:file:/home/acadgild/install/hadoop/hadoop-2.6.5/share/hadoop/common/lib/slf4j-log4j12-1.7.5.jar!/org/slf4j/impl/StaticLoggerBinder.class]

SLF4J: Found binding in [jar:file:/home/acadgild/install/hbase/hbase-1.2.6/lib/slf4j-log4j12-1.7.5.jar!/org/slf4j/impl/StaticLoggerBinder.class]

SLF4J: See [http://www.slf4j.org/codes.html#multiple\\_bindings](http://www.slf4j.org/codes.html#multiple_bindings) for an explanation.

SLF4J: Actual binding is of type [org.slf4j.impl.Log4jLoggerFactory]

18/07/07 12:05:26 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable

18/07/07 12:05:26 INFO Configuration.deprecation: mapred.jar is deprecated. Instead, use mapreduce.job.jar

18/07/07 12:05:26 INFO Configuration.deprecation: mapred.map.tasks is deprecated. Instead, use mapreduce.job.maps

18/07/07 12:05:26 INFO client.RMProxy: Connecting to ResourceManager at localhost/127.0.0.1:8032

Sat Jul 07 12:05:28 IST 2018 WARN: Establishing SSL connection without server's identity verification is not recommended. According to MySQL 5.5.45+, 5.6.26+ and 5.7.6+ requirements SSL connection must be established by default if explicit option isn't set. For compliance with existing applications not using SSL the verifyServerCertificate property is set to 'false'. You need either to explicitly disable SSL by setting useSSL=false, or set useSSL=true and provide truststore for server certificate verification.

18/07/07 12:05:28 INFO db.DBInputFormat: Using read committed transaction isolation

18/07/07 12:05:28 INFO mapreduce.JobSubmitter: number of splits:1

18/07/07 12:05:28 INFO mapreduce.JobSubmitter: Submitting tokens for job: job\_1530934123789\_0012

18/07/07 12:05:29 INFO impl.YarnClientImpl: Submitted application application\_1530934123789\_0012

18/07/07 12:05:29 INFO mapreduce.Job: The url to track the job: http://localhost:8088/proxy/application\_1530934123789\_0012/

18/07/07 12:05:29 INFO mapreduce.Job: Running job: job\_1530934123789\_0012

18/07/07 12:05:36 INFO mapreduce.Job: Job job\_1530934123789\_0012 running in uber mode : false

18/07/07 12:05:36 INFO mapreduce.Job: map 0% reduce 0%

18/07/07 12:05:42 INFO mapreduce.Job: map 100% reduce 0%

18/07/07 12:05:43 INFO mapreduce.Job: Job job\_1530934123789\_0012 completed successfully

18/07/07 12:05:43 INFO mapreduce.Job: Counters: 30

File System Counters

FILE: Number of bytes read=0

```

FILE: Number of bytes written=127863
FILE: Number of read operations=0
FILE: Number of large read operations=0
FILE: Number of write operations=0
HDFS: Number of bytes read=87
HDFS: Number of bytes written=413
HDFS: Number of read operations=4
HDFS: Number of large read operations=0
HDFS: Number of write operations=2
Job Counters
  Launched map tasks=1
  Other local map tasks=1
  Total time spent by all maps in occupied slots (ms)=3587
  Total time spent by all reduces in occupied slots (ms)=0
  Total time spent by all map tasks (ms)=3587
  Total vcore-milliseconds taken by all map tasks=3587
  Total megabyte-milliseconds taken by all map tasks=3673088
Map-Reduce Framework
  Map input records=12
  Map output records=12
  Input split bytes=87
  Spilled Records=0
  Failed Shuffles=0
  Merged Map outputs=0
  GC time elapsed (ms)=63
  CPU time spent (ms)=1720
  Physical memory (bytes) snapshot=170852352
  Virtual memory (bytes) snapshot=2086330368
  Total committed heap usage (bytes)=108003328
File Input Format Counters
  Bytes Read=0
File Output Format Counters
  Bytes Written=413
18/07/07 12:05:43 INFO mapreduce.ImportJobBase: Transferred 413 bytes in 16.2691
seconds (25.3855 bytes/sec)
18/07/07 12:05:43 INFO mapreduce.ImportJobBase: Retrieved 12 records.
[acadgild.mmisra ~]$
[acadgild.mmisra ~]$
[acadgild.mmisra ~]$
[acadgild.mmisra ~]$ hadoop fs -ls /
18/07/07 12:05:50 WARN util.NativeCodeLoader: Unable to load native-hadoop library for
your platform... using builtin-java classes where applicable
Found 8 items
drwxr-xr-x - acadgild supergroup          0 2018-07-07 11:27 /files
drwxr-xr-x - acadgild supergroup          0 2018-06-24 10:50 /mohit
drwxr-xr-x - acadgild supergroup          0 2018-07-07 11:59 /sqoopout
drwxr-xr-x - acadgild supergroup          0 2018-07-07 10:34 /sqoopoutbyid
drwxr-xr-x - acadgild supergroup          0 2018-07-07 12:05 /sqoopoutfull
-rw-r--r-- 1 acadgild supergroup        336 2018-07-07 10:54 /test.txt
drwx----- - acadgild supergroup          0 2018-07-07 11:26 /tmp
drwxr-xr-x - acadgild supergroup          0 2018-07-03 14:12 /user
[acadgild.mmisra ~]$ hadoop fs -ls /sqoopoutfull
18/07/07 12:06:04 WARN util.NativeCodeLoader: Unable to load native-hadoop library for
your platform... using builtin-java classes where applicable
Found 2 items
-rw-r--r-- 1 acadgild supergroup          0 2018-07-07 12:05 /sqoopoutfull/ SUCCESS
-rw-r--r-- 1 acadgild supergroup        413 2018-07-07 12:05 /sqoopoutfull/part-m-
00000
You have new mail in /var/spool/mail/acadgild

```

### #Dump content of the HDFS file

```
[acadgild.mmisra ~]$ hadoop fs -ls /sqoopoutfull/part-m-00000
18/07/07 12:06:12 WARN util.NativeCodeLoader: Unable to load native-hadoop library for
your platform... using builtin-java classes where applicable
-rw-r--r--  1 acadgild supergroup      413 2018-07-07 12:05 /sqoopoutfull/part-m-
00000
```

### # 12 records successfully imported in HDFS

```
[acadgild.mmisra ~]$ hadoop fs -cat /sqoopoutfull/part-m-00000
18/07/07 12:06:19 WARN util.NativeCodeLoader: Unable to load native-hadoop library for
your platform... using builtin-java classes where applicable
1,Rizvi,Syed,Whitefield,Bangalore
2,Singh,Amit,Whitefield,Bangalore
3,Singh,Ravi,Whitefield,Bangalore
4,Misra,Raju,Whitefield,Bangalore
5,hulu,bulu,Whitefield,Bangalore
6,Sinha,bulbul,Whitefield,Bangalore
7,Chue,Wang,Whitefield,Bangalore
8,Mehra,Vinod,Whitefield,Bangalore
9,Bachchan,Amitabh,Whitefield,Bombay
10,Bai,Munni,Whitefield,Lucknow
11,Bhai,Laddoo,Koramangala,Lucknow
12,Hosakote,Sri,Koramangala,Lucknow
```

### # create a new table Person2 in mysql with the same schema as table Person

```
mysql> create table Person2 like Person
```

```
-> ;
```

```
Query OK, 0 rows affected (0.02 sec)
```

```
mysql> describe table Person2;
```

```
ERROR 1064 (42000): You have an error in your SQL syntax; check the manual that
corresponds to your MySQL server version for the right syntax to use near 'table
Person2' at line 1
```

```
mysql> describe Person2;
```

Field	Type	Null	Key	Default	Extra
person_id	int(11)	NO	PRI	NULL	
lname	varchar(20)	YES		NULL	
fname	varchar(20)	YES		NULL	
area	varchar(20)	YES		NULL	
city	varchar(20)	YES		NULL	

```
5 rows in set (0.00 sec)
```

```
mysql> select * Person2;
```

```
ERROR 1064 (42000): You have an error in your SQL syntax; check the manual that
corresponds to your MySQL server version for the right syntax to use near 'Person2' at
line 1
```

### # see if the table is empty

```
mysql> select * from Person2;
```

```
Empty set (0.00 sec)
```

## #export data into this table Person2 from HDFS which was created in the previous step

```
[acadgild.mmisra ~]$sqoop export --connect jdbc:mysql://localhost/simplidb --table
Person2 --username root -password Root@123 --export-dir /sqoopoutfull/part-m-00000
Warning: /home/acadgild/install/sqoop/sqoop-1.4.6.bin__hadoop-2.0.4-alpha/./hcatalog
does not exist! HCatalog jobs will fail.
Please set $HCAT_HOME to the root of your HCatalog installation.
Warning: /home/acadgild/install/sqoop/sqoop-1.4.6.bin__hadoop-2.0.4-alpha/./accumulo
does not exist! Accumulo imports will fail.
Please set $ACCUMULO_HOME to the root of your Accumulo installation.
18/07/07 12:07:08 INFO sqoop.Sqoop: Running Sqoop version: 1.4.6
18/07/07 12:07:08 WARN tool.BaseSqoopTool: Setting your password on the command-line
is insecure. Consider using -P instead.
18/07/07 12:07:08 INFO manager.MySQLManager: Preparing to use a MySQL streaming
resultset.
18/07/07 12:07:08 INFO tool.CodeGenTool: Beginning code generation
Sat Jul 07 12:07:08 IST 2018 WARN: Establishing SSL connection without server's
identity verification is not recommended. According to MySQL 5.5.45+, 5.6.26+ and
5.7.6+ requirements SSL connection must be established by default if explicit option
isn't set. For compliance with existing applications not using SSL the
verifyServerCertificate property is set to 'false'. You need either to explicitly
disable SSL by setting useSSL=false, or set useSSL=true and provide truststore for
server certificate verification.
18/07/07 12:07:09 INFO manager.SqlManager: Executing SQL statement: SELECT t.* FROM
`Person2` AS t LIMIT 1
18/07/07 12:07:09 INFO manager.SqlManager: Executing SQL statement: SELECT t.* FROM
`Person2` AS t LIMIT 1
18/07/07 12:07:09 INFO orm.CompilationManager: HADOOP_MAPRED_HOME is
/home/acadgild/install/hadoop/hadoop-2.6.5
Note: /tmp/sqoop-acadgild/compile/2ec2c4af530d8064588c7c322703f80a/Person2.java uses
or overrides a deprecated API.
Note: Recompile with -Xlint:deprecation for details.
18/07/07 12:07:11 INFO orm.CompilationManager: Writing jar file: /tmp/sqoop-
acadgild/compile/2ec2c4af530d8064588c7c322703f80a/Person2.jar
18/07/07 12:07:11 INFO mapreduce.ExportJobBase: Beginning export of Person2
SLF4J: Class path contains multiple SLF4J bindings.
SLF4J: Found binding in [jar:file:/home/acadgild/install/hadoop/hadoop-
2.6.5/share/hadoop/common/lib/slf4j-log4j12-
1.7.5.jar!/org/slf4j/impl/StaticLoggerBinder.class]
SLF4J: Found binding in [jar:file:/home/acadgild/install/hbase/hbase-1.2.6/lib/slf4j-
log4j12-1.7.5.jar!/org/slf4j/impl/StaticLoggerBinder.class]
SLF4J: See http://www.slf4j.org/codes.html#multiple_bindings for an explanation.
SLF4J: Actual binding is of type [org.slf4j.impl.Log4jLoggerFactory]
18/07/07 12:07:11 WARN util.NativeCodeLoader: Unable to load native-hadoop library for
your platform... using builtin-java classes where applicable
18/07/07 12:07:11 INFO Configuration.deprecation: mapred.jar is deprecated. Instead,
use mapreduce.job.jar
18/07/07 12:07:12 INFO Configuration.deprecation:
mapred.reduce.tasks.speculative.execution is deprecated. Instead, use
mapreduce.reduce.speculative
18/07/07 12:07:12 INFO Configuration.deprecation:
mapred.map.tasks.speculative.execution is deprecated. Instead, use
mapreduce.map.speculative
18/07/07 12:07:12 INFO Configuration.deprecation: mapred.map.tasks is deprecated.
Instead, use mapreduce.job.maps
18/07/07 12:07:12 INFO client.RMProxy: Connecting to ResourceManager at
localhost/127.0.0.1:8032
18/07/07 12:07:13 INFO input.FileInputFormat: Total input paths to process : 1
18/07/07 12:07:13 INFO input.FileInputFormat: Total input paths to process : 1
18/07/07 12:07:13 INFO mapreduce.JobSubmitter: number of splits:4
18/07/07 12:07:13 INFO Configuration.deprecation:
mapred.map.tasks.speculative.execution is deprecated. Instead, use
mapreduce.map.speculative
```

```
18/07/07 12:07:14 INFO mapreduce.JobSubmitter: Submitting tokens for job:
job_1530934123789_0013
18/07/07 12:07:14 INFO impl.YarnClientImpl: Submitted application
application_1530934123789_0013
18/07/07 12:07:14 INFO mapreduce.Job: The url to track the job:
http://localhost:8088/proxy/application_1530934123789_0013/
18/07/07 12:07:14 INFO mapreduce.Job: Running job: job_1530934123789_0013
18/07/07 12:07:22 INFO mapreduce.Job: Job job_1530934123789_0013 running in uber mode
: false
18/07/07 12:07:22 INFO mapreduce.Job: map 0% reduce 0%
18/07/07 12:07:37 INFO mapreduce.Job: map 25% reduce 0%
18/07/07 12:07:38 INFO mapreduce.Job: map 100% reduce 0%
18/07/07 12:07:38 INFO mapreduce.Job: Job job_1530934123789_0013 completed
successfully
18/07/07 12:07:38 INFO mapreduce.Job: Counters: 30
    File System Counters
        FILE: Number of bytes read=0
        FILE: Number of bytes written=510228
        FILE: Number of read operations=0
        FILE: Number of large read operations=0
        FILE: Number of write operations=0
        HDFS: Number of bytes read=1677
        HDFS: Number of bytes written=0
        HDFS: Number of read operations=19
        HDFS: Number of large read operations=0
        HDFS: Number of write operations=0
    Job Counters
        Launched map tasks=4
        Data-local map tasks=4
        Total time spent by all maps in occupied slots (ms)=54000
        Total time spent by all reduces in occupied slots (ms)=0
        Total time spent by all map tasks (ms)=54000
        Total vcore-milliseconds taken by all map tasks=54000
        Total megabyte-milliseconds taken by all map tasks=55296000
    Map-Reduce Framework
        Map input records=12
        Map output records=12
        Input split bytes=576
        Spilled Records=0
        Failed Shuffles=0
        Merged Map outputs=0
        GC time elapsed (ms)=708
        CPU time spent (ms)=7090
        Physical memory (bytes) snapshot=637562880
        Virtual memory (bytes) snapshot=8342564864
        Total committed heap usage (bytes)=451936256
    File Input Format Counters
        Bytes Read=0
    File Output Format Counters
        Bytes Written=0
18/07/07 12:07:38 INFO mapreduce.ExportJobBase: Transferred 1.6377 KB in 26.1849
seconds (64.0445 bytes/sec)
18/07/07 12:07:38 INFO mapreduce.ExportJobBase: Exported 12 records.
You have new mail in /var/spool/mail/acadgild
[acadgild.mmisra ~]$
```

## #Check contents of table Person2 in MySql , there should be 12 records

```
mysql> select * from Person2;
```

person_id	lname	fname	area	city
1	Rizvi	Syed	Whitefield	Bangalore
2	Singh	Amit	Whitefield	Bangalore
3	Singh	Ravi	Whitefield	Bangalore
4	Misra	Raju	Whitefield	Bangalore
5	hulu	bulu	Whitefield	Bangalore
6	Sinha	bulbul	Whitefield	Bangalore
7	Chue	Wang	Whitefield	Bangalore
8	Mehra	Vinod	Whitefield	Bangalore
9	Bachchan	Amitabh	Whitefield	Bombay
10	Bai	Munni	Whitefield	Lucknow
11	Bhai	Laddoo	Koramangala	Lucknow
12	Hosakote	Sri	Koramangala	Lucknow

```
12 rows in set (0.00 sec)
```

```
mysql>
```

## SOLUTION – TASK2

### #import data from table Person where person\_id=3

```
[acadgild.mmisra ~]$
```

```
[acadgild.mmisra ~]$ sqoop import --connect jdbc:mysql://localhost/simplidb --table  
Person --username root --password Root@123 --target-dir /sqoopout -m 1 --where  
person_id=3
```

```
Warning: /home/acadgild/install/sqoop/sqoop-1.4.6.bin__hadoop-2.0.4-alpha/../../hcatalog  
does not exist! HCatalog jobs will fail.
```

```
Please set $HCAT_HOME to the root of your HCatalog installation.
```

```
Warning: /home/acadgild/install/sqoop/sqoop-1.4.6.bin__hadoop-2.0.4-alpha/../../accumulo  
does not exist! Accumulo imports will fail.
```

```
Please set $ACCUMULO_HOME to the root of your Accumulo installation.
```

```
18/07/07 11:59:31 INFO sqoop.Sqoop: Running Sqoop version: 1.4.6
```

```
18/07/07 11:59:31 WARN tool.BaseSqoopTool: Setting your password on the command-line  
is insecure. Consider using -P instead.
```

```
18/07/07 11:59:31 INFO manager.MySQLManager: Preparing to use a MySQL streaming  
resultset.
```

```
18/07/07 11:59:31 INFO tool.CodeGenTool: Beginning code generation
```

```
Sat Jul 07 11:59:31 IST 2018 WARN: Establishing SSL connection without server's  
identity verification is not recommended. According to MySQL 5.5.45+, 5.6.26+ and  
5.7.6+ablsished by default if explicit option isn't set. For compliance with existing  
applications not using SSL the verifyServerCertificate property is set to 'false'. You  
neetting useSSL=false, or set useSSL=true and provide truststore for server  
certificate verification.
```

```
18/07/07 11:59:32 INFO manager.SqlManager: Executing SQL statement: SELECT t.* FROM  
`Person` AS t LIMIT 1
```

```
18/07/07 11:59:32 INFO manager.SqlManager: Executing SQL statement: SELECT t.* FROM  
`Person` AS t LIMIT 1
```

```
18/07/07 11:59:32 INFO orm.CompilationManager: HADOOP_MAPRED_HOME is  
/home/acadgild/install/hadoop/hadoop-2.6.5
```

```
Note: /tmp/sqoop-acadgild/compile/495276af6e042934a57d226c53226bf6/Person.java uses or  
overrides a deprecated API.
```

```
Note: Recompile with -Xlint:deprecation for details.
```

```
18/07/07 11:59:34 INFO orm.CompilationManager: Writing jar file: /tmp/sqoop-  
acadgild/compile/495276af6e042934a57d226c53226bf6/Person.jar
```

```
18/07/07 11:59:34 WARN manager.MySQLManager: It looks like you are importing from
```



```

mysql.
18/07/07 11:59:34 WARN manager.MySQLManager: This transfer can be faster! Use the --
direct
18/07/07 11:59:34 WARN manager.MySQLManager: option to exercise a MySQL-specific fast
path.
18/07/07 11:59:34 INFO manager.MySQLManager: Setting zero DATETIME behavior to
convertToNull (mysql)
18/07/07 11:59:34 INFO mapreduce.ImportJobBase: Beginning import of Person
SLF4J: Class path contains multiple SLF4J bindings.
SLF4J: Found binding in [jar:file:/home/acadgild/install/hadoop/hadoop-
2.6.5/share/hadoop/common/lib/slf4j-log4j12-
1.7.5.jar!/org/slf4j/impl/StaticLoggerBinder.class]
SLF4J: Found binding in [jar:file:/home/acadgild/install/hbase/hbase-1.2.6/lib/slf4j-
log4j12-1.7.5.jar!/org/slf4j/impl/StaticLoggerBinder.class]
SLF4J: See http://www.slf4j.org/codes.html#multiple_bindings for an explanation.
SLF4J: Actual binding is of type [org.slf4j.impl.Log4jLoggerFactory]
18/07/07 11:59:34 WARN util.NativeCodeLoader: Unable to load native-hadoop library for
your platform... using builtin-java classes where applicable
18/07/07 11:59:34 INFO Configuration.deprecation: mapred.jar is deprecated. Instead,
use mapreduce.job.jar
18/07/07 11:59:35 INFO Configuration.deprecation: mapred.map.tasks is deprecated.
Instead, use mapreduce.job.maps
18/07/07 11:59:35 INFO client.RMProxy: Connecting to ResourceManager at
localhost/127.0.0.1:8032
Sat Jul 07 11:59:36 IST 2018 WARN: Establishing SSL connection without server's
identity verification is not recommended. According to MySQL 5.5.45+, 5.6.26+ and
5.7.6+ablsihed by default if explicit option isn't set. For compliance with existing
applications not using SSL the verifyServerCertificate property is set to 'false'. You
neetting useSSL=false, or set useSSL=true and provide truststore for server
certificate verification.
18/07/07 11:59:36 INFO db.DBInputFormat: Using read committed transaction isolation
18/07/07 11:59:36 INFO mapreduce.JobSubmitter: number of splits:1
18/07/07 11:59:36 INFO mapreduce.JobSubmitter: Submitting tokens for job:
job_1530934123789_0010
18/07/07 11:59:37 INFO impl.YarnClientImpl: Submitted application
application_1530934123789_0010
18/07/07 11:59:37 INFO mapreduce.Job: The url to track the job:
http://localhost:8088/proxy/application_1530934123789_0010/
18/07/07 11:59:37 INFO mapreduce.Job: Running job: job_1530934123789_0010
18/07/07 11:59:45 INFO mapreduce.Job: Job job_1530934123789_0010 running in uber mode
: false
18/07/07 11:59:45 INFO mapreduce.Job: map 0% reduce 0%
18/07/07 11:59:52 INFO mapreduce.Job: map 100% reduce 0%
18/07/07 11:59:52 INFO mapreduce.Job: Job job_1530934123789_0010 completed
successfully
18/07/07 11:59:52 INFO mapreduce.Job: Counters: 30
    File System Counters
        FILE: Number of bytes read=0
        FILE: Number of bytes written=128011
        FILE: Number of read operations=0
        FILE: Number of large read operations=0
        FILE: Number of write operations=0
        HDFS: Number of bytes read=87
        HDFS: Number of bytes written=34
        HDFS: Number of read operations=4
        HDFS: Number of large read operations=0
        HDFS: Number of write operations=2
    Job Counters
        Launched map tasks=1
        Other local map tasks=1
        Total time spent by all maps in occupied slots (ms)=4042
        Total time spent by all reduces in occupied slots (ms)=0
        Total time spent by all map tasks (ms)=4042

```

```

        Total vcore-milliseconds taken by all map tasks=4042
        Total megabyte-milliseconds taken by all map tasks=4139008
Map-Reduce Framework
    Map input records=1
    Map output records=1
    Input split bytes=87
    Spilled Records=0
    Failed Shuffles=0
    Merged Map outputs=0
    GC time elapsed (ms)=87
    CPU time spent (ms)=1780
    Physical memory (bytes) snapshot=170639360
    Virtual memory (bytes) snapshot=2086248448
    Total committed heap usage (bytes)=116916224
File Input Format Counters
    Bytes Read=0
File Output Format Counters
    Bytes Written=34
18/07/07 11:59:52 INFO mapreduce.ImportJobBase: Transferred 34 bytes in 17.3544
seconds (1.9592 bytes/sec)
18/07/07 11:59:52 INFO mapreduce.ImportJobBase: Retrieved 1 records.

```

### #Create a new table person1 with same schema as table Person

```

mysql> create table Person1 like Person;
Query OK, 0 rows affected (0.03 sec)

```

```

mysql> show tables;
+-----+
| Tables_in_simplidb |
+-----+
| Person              |
| Person1             |
| employee            |
+-----+
3 rows in set (0.00 sec)

```

```

mysql> describe Person1;
+-----+-----+-----+-----+-----+-----+
| Field      | Type          | Null | Key | Default | Extra |
+-----+-----+-----+-----+-----+-----+
| person_id | int(11)       | NO   | PRI | NULL    |       |
| lname     | varchar(20)   | YES  |     | NULL    |       |
| fname     | varchar(20)   | YES  |     | NULL    |       |
| area      | varchar(20)   | YES  |     | NULL    |       |
| city      | varchar(20)   | YES  |     | NULL    |       |
+-----+-----+-----+-----+-----+-----+
5 rows in set (0.01 sec)

```

```
mysql>
```

### #Export data to the Person1 table from HDFS which was created in previous step for person\_id=3

```

[acadgild.mmisra ~]$sqoop export --connect jdbc:mysql://localhost/simplidb --table
Person1 --username root -password Root@123 --export-dir /sqoopout/part-m-00000
Warning: /home/acadgild/install/sqoop/sqoop-1.4.6.bin__hadoop-2.0.4-alpha/./hcatalog
does not exist! HCatalog jobs will fail.
Please set $HCAT_HOME to the root of your HCatalog installation.
Warning: /home/acadgild/install/sqoop/sqoop-1.4.6.bin__hadoop-2.0.4-alpha/./accumulo
does not exist! Accumulo imports will fail.

```

Please set \$ACCUMULO\_HOME to the root of your Accumulo installation.

18/07/07 12:02:35 INFO sqoop.Sqoop: Running Sqoop version: 1.4.6

18/07/07 12:02:36 WARN tool.BaseSqoopTool: Setting your password on the command-line is insecure. Consider using -P instead.

18/07/07 12:02:36 INFO manager.MySQLManager: Preparing to use a MySQL streaming resultset.

18/07/07 12:02:36 INFO tool.CodeGenTool: Beginning code generation

Sat Jul 07 12:02:36 IST 2018 WARN: Establishing SSL connection without server's identity verification is not recommended. According to MySQL 5.5.45+, 5.6.26+ and 5.7.6+ requirements SSL connection must be established by default if explicit option isn't set. For compliance with existing applications not using SSL the verifyServerCertificate property is set to 'false'. You need either to explicitly disable SSL by setting useSSL=false, or set useSSL=true and provide truststore for server certificate verification.

18/07/07 12:02:36 INFO manager.SqlManager: Executing SQL statement: SELECT t.\* FROM `Person1` AS t LIMIT 1

18/07/07 12:02:36 INFO manager.SqlManager: Executing SQL statement: SELECT t.\* FROM `Person1` AS t LIMIT 1

18/07/07 12:02:36 INFO orm.CompilationManager: HADOOP\_MAPRED\_HOME is /home/acadgild/install/hadoop/hadoop-2.6.5

Note: /tmp/sqoop-acadgild/compile/548df9a3ffbc294327c1e0ed96b43b13/Person1.java uses or overrides a deprecated API.

Note: Recompile with -Xlint:deprecation for details.

18/07/07 12:02:38 INFO orm.CompilationManager: Writing jar file: /tmp/sqoop-acadgild/compile/548df9a3ffbc294327c1e0ed96b43b13/Person1.jar

18/07/07 12:02:38 INFO mapreduce.ExportJobBase: Beginning export of Person1

SLF4J: Class path contains multiple SLF4J bindings.

SLF4J: Found binding in [jar:file:/home/acadgild/install/hadoop/hadoop-2.6.5/share/hadoop/common/lib/slf4j-log4j12-1.7.5.jar!/org/slf4j/impl/StaticLoggerBinder.class]

SLF4J: Found binding in [jar:file:/home/acadgild/install/hbase/hbase-1.2.6/lib/slf4j-log4j12-1.7.5.jar!/org/slf4j/impl/StaticLoggerBinder.class]

SLF4J: See [http://www.slf4j.org/codes.html#multiple\\_bindings](http://www.slf4j.org/codes.html#multiple_bindings) for an explanation.

SLF4J: Actual binding is of type [org.slf4j.impl.Log4jLoggerFactory]

18/07/07 12:02:39 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable

18/07/07 12:02:39 INFO Configuration.deprecation: mapred.jar is deprecated. Instead, use mapreduce.job.jar

18/07/07 12:02:40 INFO Configuration.deprecation: mapred.reduce.tasks.speculative.execution is deprecated. Instead, use mapreduce.reduce.speculative

18/07/07 12:02:40 INFO Configuration.deprecation: mapred.map.tasks.speculative.execution is deprecated. Instead, use mapreduce.map.speculative

18/07/07 12:02:40 INFO Configuration.deprecation: mapred.map.tasks is deprecated. Instead, use mapreduce.job.maps

18/07/07 12:02:40 INFO client.RMProxy: Connecting to ResourceManager at localhost/127.0.0.1:8032

18/07/07 12:02:42 INFO input.FileInputFormat: Total input paths to process : 1

18/07/07 12:02:42 INFO input.FileInputFormat: Total input paths to process : 1

18/07/07 12:02:42 INFO mapreduce.JobSubmitter: number of splits:4

18/07/07 12:02:42 INFO Configuration.deprecation: mapred.map.tasks.speculative.execution is deprecated. Instead, use mapreduce.map.speculative

18/07/07 12:02:42 INFO mapreduce.JobSubmitter: Submitting tokens for job: job\_1530934123789\_0011

18/07/07 12:02:42 INFO impl.YarnClientImpl: Submitted application application\_1530934123789\_0011

18/07/07 12:02:43 INFO mapreduce.Job: The url to track the job: [http://localhost:8088/proxy/application\\_1530934123789\\_0011/](http://localhost:8088/proxy/application_1530934123789_0011/)

18/07/07 12:02:43 INFO mapreduce.Job: Running job: job\_1530934123789\_0011

18/07/07 12:02:51 INFO mapreduce.Job: Job job\_1530934123789\_0011 running in uber mode : false

```

18/07/07 12:02:51 INFO mapreduce.Job: map 0% reduce 0%
18/07/07 12:03:07 INFO mapreduce.Job: map 100% reduce 0%
18/07/07 12:03:08 INFO mapreduce.Job: Job job_1530934123789_0011 completed
successfully
18/07/07 12:03:08 INFO mapreduce.Job: Counters: 30
    File System Counters
        FILE: Number of bytes read=0
        FILE: Number of bytes written=510212
        FILE: Number of read operations=0
        FILE: Number of large read operations=0
        FILE: Number of write operations=0
        HDFS: Number of bytes read=664
        HDFS: Number of bytes written=0
        HDFS: Number of read operations=19
        HDFS: Number of large read operations=0
        HDFS: Number of write operations=0
    Job Counters
        Launched map tasks=4
        Data-local map tasks=4
        Total time spent by all maps in occupied slots (ms)=56522
        Total time spent by all reduces in occupied slots (ms)=0
        Total time spent by all map tasks (ms)=56522
        Total vcore-milliseconds taken by all map tasks=56522
        Total megabyte-milliseconds taken by all map tasks=57878528
    Map-Reduce Framework
        Map input records=1
        Map output records=1
        Input split bytes=556
        Spilled Records=0
        Failed Shuffles=0
        Merged Map outputs=0
        GC time elapsed (ms)=1079
        CPU time spent (ms)=7600
        Physical memory (bytes) snapshot=673996800
        Virtual memory (bytes) snapshot=8339111936
        Total committed heap usage (bytes)=485490688
    File Input Format Counters
        Bytes Read=0
    File Output Format Counters
        Bytes Written=0
18/07/07 12:03:08 INFO mapreduce.ExportJobBase: Transferred 664 bytes in 28.5255
seconds (23.2774 bytes/sec)
18/07/07 12:03:08 INFO mapreduce.ExportJobBase: Exported 1 records.
You have new mail in /var/spool/mail/acadgild

```

### #dump contents of table Person1 to verify successful export

```

mysql> select * from Person1;
+-----+-----+-----+-----+-----+
| person_id | lname | fname | area          | city          |
+-----+-----+-----+-----+-----+
|          3 | Singh | Ravi  | Whitefield    | Bangalore     |
+-----+-----+-----+-----+-----+
1 row in set (0.00 sec)

```